

# Bioinformatics Hub Phoenix cluster

## Users

### Bioinformatics Hub

Stephen Pederson

Hien To

Jimmy Breen (a1650598)

Alastair Ludington (a)

Lakshmi Chandrapaty (a1690027)

### Robinson Research Institute

Jimmy Breen (a1650598)

Ben Mayne

Shalem Leemaqz

## Setup

To login to Phoenix via ssh. You will need your university account id

ssh a1650598@phoenix.rc.adelaide.edu.au

It will ask you for a password. This is your university account password.

You dont really want to have keep on typing that everytime you login, so lets setup an ssh config file. If you've ever used ssh, you will have an .ssh directory in your home folder

```
cd ~/.ssh/
```

This contains all your ssh relevant stuff like ssh keypairs (will explain later) and authorized keys. To create an alias for the phoenix cluster, lets create the config file (if you dont have one already)

```
echo "Host phoenix
HostName phoenix.rc.adelaide.edu.au
User a1650598
ServerAliveInterval 60" >> .ssh/config
```

Now you can just login to Phoenix, by doing the following:

```
ssh phoenix
```

You still need to put in your password. But we can skip this step easily by using a key-pair. A key-pair is like a handshake between two networks (this being your local computer and the Phoenix cluster) that enables password-less access to the desired network. You can do this between any two networks that you can access via ssh.

The steps i'll run through now are contained in [this link](#)

### Step One—Create the RSA Key Pair

The first step is to create the key pair on the client machine (there is a good chance that this will just be your computer):

```
ssh-keygen -t rsa
```

It will ask you for a pass-phrase (THIS IS VERY IMPORTANT TO REMEMBER...but then again, all passwords are fairly important i guess). We will use RSA encryption here, but there are others that you could potentially try (dsa for example). I would suggest using this for now

```
Enter passphrase (empty for no passphrase):
```

The entire process will probably look like this:

```
ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/demo/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/demo/.ssh/id_rsa.
Your public key has been saved in /home/demo/.ssh/id_rsa.pub.
The key fingerprint is:
4a:dd:0a:c6:35:4e:3f:ed:27:38:8c:74:44:4d:93:67 demo@a
The key's randomart image is:
+--[ RSA 2048 ]-----+
|           .oo.      |
|          .  o.E     |
|         + .  o      |
|        . = = .      |
|       = S = .       |
|      o + = +        |
|     . o + o .       |
|          . o        |
|                     |
+-----+
```

What has been made are two things:

1. A encrypted key; and 2. A public key that needs to be transfered to the *other* server that

you want to talk to

For me, this would be:

```
ssh-copy-id a1650598@phoenix.rc.adelaide.edu.au
```

```
The authenticity of host '12.34.56.78 (12.34.56.78)' can't be established.  
RSA key fingerprint is b1:2d:33:67:ce:35:4d:5f:f3:a8:cd:c0:c4:48:86:12.  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added '12.34.56.78' (RSA) to the list of known hosts.  
user@12.34.56.78's password:
```

The keys should be in a "public key" file called:

```
~/.ssh/authorized_keys
```

Try ssh'ing in and seeing if everything went through

SSH keys are obviously very sensitive, so they won't work if the file permissions are not set to read-only. We don't want anyone editing them or deleting them.

## Introduction

Ok we're in, now what do we do with the cluster. First, let's have a look at the file-system that's attached

```
[a1650598@phoenix1 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
rootfs          99G   36G   58G   39% /
10.33.139.1:/opt/boot/hpc6
                99G   36G   58G   39% /
none            24G  140K   24G    1% /dev
/dev/ram        10G   29M   10G    1% /ram
/dev/sda1       7.8G   4.7G   2.7G   64% /localscratch
tmpfs           24G   12K   24G    1% /dev/shm
phoenixnfs1.rc.adelaide.edu.au:/home
                40T   7.7T   33T   20% /home
phoenixnfs1.rc.adelaide.edu.au:/opt/shared
                1.0T  148G   876G   15% /opt/shared
phoenixnfs1.rc.adelaide.edu.au:/opt/shared/slurm/slurm.state
                1.0T  148G   876G   15% /opt/shared/slurm/slurm.state
10.33.139.231@tcp:10.33.139.230@tcp:/data
                349T  133T  199T   40% /data
```

The two areas which you will spend your time is:

- /home - This is where you put your custom programs, settings and store some data
- /localscratch: Scratch storage is the fastest storage, so ideally you want to run all your major analyses here

When you arrive in the machine, you should create a directory on localscratch. For example:

```
mkdir /localscratch/jbreen/
```

## Modules & Programs

There are two types of programs on phoenix, ones that have been installed by us (i.e. me) locally and "modules"

Modules are convenient little commands that allow you to run a simple command to use important programs.

For example, If I wanted to use bowtie to run an alignment, but I don't have that program in my PATH, I can do the following:

```
[a1650598@phoenix11 Programs]$ module avail
```

```
-----  
/usr/share/Modules/modulefiles -----  
-----
```

```
dot          module-git  module-info modules      null          use.own
```

```
-----  
/etc/modulefiles -----  
-----
```

OpenFOAM/2.1.1	cmake/2.8.12.1	gslib
mpfr/3.1.2	pythia/8.2.10	
OpenFOAM/2.3.1	cuda/5.0	gurobi/6.0.5
muscle/3.8.31	python/2.7.10	
R/3.2.0	cuda/6.0	hdf5/1.8.15
namd/2.10-cuda	qt/4.8.6	
abaqus/6.14	cuda/6.5	hdf5/1.8.15-patch1
ncbi-blast/2.2.24+	root/5.34.32	
ansys/14.5	cuda/7.0	hkl2000/1
ncbi-blast/2.2.27+	samtools/1.2	
ansys/16.1	cufflinks/2.2.1	hoomd-blue/1.1.1
ncbi-blast/2.2.31+	scotch/6.0.0	
ansys/16.2	epmr/15.04	hoomd-blue/1.2.0
netcdf/4.4.0-C	subread/1.4.6-p2	
atlas/3.10.2	fastqc/0.11.3	hoomd-blue/hoomd-blue
netcdf/4.4.2-Fortran	terachem/v1.50k	
bedtools/2.25.0	fastx_toolkit/0.0.14	intel/13.1.3.174
netcdf/4.4.2-Fortran-intel	tophat/2.1.0	
bioperl/1.6.924	flex/2.5.39	intelmpi/5.0.3.048
opencv/2.4.10	trim_galore/0.4.0	
blast/2.2.24	gaussian/g09	java/1.8.0
openmm/6.2.0	trinity/2.0.6	
boost/1.54.0	gcc/4.8.4	julia/0.3
openmm/6.2.0-cuda	uofa_utils/1	
boost/1.55.0	gcc/5.1.0	lammps/2015.05.15
openmpi/gnu/1.8.4	wamit/v7101	
bowtie/1.1.1	glibc/2.16.0	lammps/2015.08.10/gpu
openmpi/intel/1.8.1	wu-blast/2.0	
bowtie/2.2.5	glog/0.3.3	lhpdf/5.9.1
openmpi-x86_64	xds/1	
bwa/0.7.12	gmp/5.1.2	ls-dyna/8.0.0
paraview/4.1.0	zlib/1.2.8	
ccp4/6.5	gnu-parallel/20150322	matlab/2014a
pgi/15.4		
sensor/4.2.29	go/1.5	matlab/2015a
phenix/1.9-1692		
cgal/4.3	gsl/1.9	mpc/1.0.1
protobuf/2.6.1		

I see the module for bowtie is listed there on the available modules. Now I want to load it:

```
[a1650598@phoenix11 Programs]$ module load bowtie/2.2.5
```

Lets see if I can run the program:

```
[a1650598@phoenix11 Programs]$ bowtie2
No index, query, or output file specified!
Bowtie 2 version 2.2.5 by Ben Langmead (langmea@cs.jhu.edu,
www.cs.jhu.edu/~langmea)
Usage:
  bowtie2 [options]* -x <bt2-idx> {-1 <m1> -2 <m2> | -U <r>} [-S <sam>]

<bt2-idx>  Index filename prefix (minus trailing .X.bt2).
          NOTE: Bowtie 1 and Bowtie 2 indexes are not compatible.
```

So thats modules, but you might find that you dont have module command for a program that you need. This will happen quite often, so its best to locally install your program. Generally you will need the admins (Stephen, Hien or Jimmy) to do this step because it requires a admin password.

For programs that have already been installed, have a look in the directory  
/localscratch/Programs/

```
[a1650598@phoenix11 Programs]$ ls -l /localscratch/Programs/
total 196968
drwxr-xr-x  8 a1650598      4096 Jun 11  2015 AdmixTools-20150610-git-3065acc5
-rwxr-xr-x  1 a1650598 144533833 Jul 24  2015 AdmixTools.tar.gz
drwxr-xr-x  2 a1650598      4096 Jun  2  2015 Eigensoft_6_0_1
drwxr-xr-x  6 a1650598      4096 Feb 17  09:40 GBSX
drwxr-sr-x  7 a1650598      4096 Feb  4  21:27 HiC-Pro
drwxr-sr-x 10 a1650598      4096 Feb 13  10:46 HiCExplorer
-rwxr-xr-x  1 a1650598    35147 Nov 21  2013 LICENSE
drwxr-xr-x  2 a1650598      4096 Jun  5  2015 MafFilter
drwxr-sr-x  4 a1650598      4096 Sep  3  2015 Metaxa2_2.0.2
drwxr-sr-x  5 a1650598      4096 Aug 31  2015 PileOMeth
drwxr-sr-x 12 a1650598      4096 Jun  1  2015 PyVCF
drwxr-xr-x  6 a1650598      4096 Jun 11  2015 RAxML-git20150511-297bc55146
-rwxr-xr-x  1 a1650598      280 Sep  1  2015 README.afw-acad2-rsync
drwxr-sr-x 155 a1650598      4096 Feb 28  10:41 R_libs
drwxr-sr-x  4 a1650598      4096 Dec 19  04:09 SailfishBeta-0.9.0_DebianSqueeze
drwxr-sr-x  4 a1650598      4096 Jan 13  06:31 SalmonBeta-0.6.1_DebianSqueeze
drwxr-xr-x  8 a1650598      4096 Jun  2  2015 adapterremoval2-git-0afe24-20150601
```

drwxr-xr-x	9	a1650598	12288	Aug	31	2015	angsd-20150831-git~8b89ba42
drwxr-xr-x	9	a1650598	12288	Aug	31	2015	angsd_0.902
-rwxr-xr-x	1	a1650598	144	Aug	31	2015	angsd_0.902_install.txt
drwxr-sr-x	9	a1650598	4096	Feb	14	13:06	autoconf-2.69
-rwxr-xr-x	1	a1650598	1927468	Apr	25	2012	autoconf-latest.tar.gz
drwxr-sr-x	4	a1650598	4096	Jun	6	2015	bamaddrg
drwxr-xr-x	6	a1650598	4096	Jun	2	2015	bamtools
drwxr-sr-x	7	a1650598	4096	Nov	17	10:45	bbmap
drwxr-xr-x	7	a1650598	4096	Jun	30	2015	bcftools
drwxr-xr-x	4	a1650598	4096	Jun	2	2015	bcftools-1.2
drwxr-xr-x	2	a1650598	4096	Sep	25	08:12	bedops_bin
-rwxr-xr-x	1	a1650598	5926099	Apr	22	2015	bedops_linux_x86_64-v2.4.14.tar.bz2
drwxr-xr-x	8	a1650598	4096	Jun	11	2015	binutils-2.23.2
-rwxr-xr-x	1	a1650598	21440347	Mar	27	2013	binutils-2.23.2.tar.bz2
drwxr-xr-x	4	a1650598	4096	Jun	2	2015	bioawk
drwxr-xr-x	3	a1650598	4096	Jun	3	2015	bycatch
drwxr-xr-x	10	a1650598	4096	Sep	1	2015	dmd-2.067.0
drwxr-xr-x	3	a1650598	4096	Jul	1	2015	drift_test
drwxr-xr-x	6	a1650598	4096	Jul	29	2015	fastStructure-20150729-gite47212f8
drwxr-sr-x	3	a1650598	4096	Sep	8	2015	gammytools
drwxr-xr-x	3	a1650598	4096	Sep	1	2015	gatk-3.3-0
drwxr-xr-x	4	a1650598	4096	Jun	2	2015	ghostscript-9.15
-rwxr-xr-x	1	a1650598	120234	Aug	17	2015	gmon.out
drwxr-xr-x	5	a1650598	4096	Jun	2	2015	gnuplot-4.6.6
drwxr-xr-x	4	a1650598	4096	Aug	25	2015	grg-utils
drwxr-xr-x	4	a1650598	4096	Sep	1	2015	heterozygosity_DoNotShare
drwxr-sr-x	3	a1650598	4096	Feb	4	21:03	hiCPro
drwxr-sr-x	6	a1650598	4096	Feb	13	10:30	hicup_v0.5.8
drwxr-sr-x	5	a1650598	4096	Apr	18	2015	hisat-0.1.6-beta
-rwxr-xr-x	1	a1650598	27352793	Apr	18	2015	hisat-0.1.6-beta-Linux_x86_64.zip
drwxr-sr-x	5	a1650598	4096	Nov	20	07:32	hisat2-2.0.1-beta
drwxr-sr-x	7	a1650598	4096	Feb	27	13:04	homer
drwxr-xr-x	6	a1650598	4096	Jun	30	2015	htslib
drwxr-xr-x	6	a1650598	4096	Jun	2	2015	htslib-1.2.1
drwxr-sr-x	3	a1650598	4096	Oct	30	04:38	kallisto_linux-v0.42.4
drwxr-xr-x	5	a1650598	4096	Jun	2	2015	libevent-2.0.22
drwxr-xr-x	16	a1650598	4096	Aug	25	2015	libgd-gd-2.1.1
drwxr-sr-x	5	a1650598	4096	Sep	1	2015	malt
drwxr-sr-x	7	a1650598	4096	Sep	19	2015	methclone
drwxr-sr-x	5	a1650598	4096	Sep	19	2015	methpipe
drwxr-sr-x	9	a1650598	4096	Sep	28	17:14	methtuple
drwxr-xr-x	7	a1650598	4096	Aug	7	2015	mLRho
drwxr-xr-x	4	a1650598	4096	Jun	2	2015	msmc-20150413
drwxr-xr-x	3	a1650598	4096	Jun	2	2015	mugsy_x86-64-v1r2.2
drwxr-xr-x	7	a1650598	4096	Jun	2	2015	ncurses-5.9
drwxr-sr-x	7	a1650598	4096	Jan	5	15:50	ngi_visualizations
drwxr-xr-x	5	a1650598	4096	Aug	31	2015	ngsPopGen-20150831-git~abeabb73

drwxr-xr-x	5	a1650598	4096	Jul	16	2015	nmap-6.49BETA4
drwxr-xr-x	3	a1650598	4096	Jun	3	2015	nuclearratescurve
drwxr-xr-x	11	a1650598	4096	Jun	9	2015	paleomix-git-2fc2560-20150601
drwxr-xr-x	10	a1650598	4096	Jun	2	2015	paleomix-git-2fc2560-20150601-nope
drwxr-sr-x	3	a1650598	4096	Aug	24	2015	perl5
drwxr-xr-x	2	a1650598	4096	Jun	9	2015	picard-1.133
drwxr-sr-x	2	a1650598	4096	Jan	18	13:02	plink-1.90
drwxr-sr-x	2	a1650598	4096	Jul	12	2014	plinkseq-0.10
drwxr-sr-x	6	a1650598	4096	Feb	14	13:07	pullseq
drwxr-sr-x	5	a1650598	4096	Sep	9	2015	rnammer
drwxr-xr-x	8	a1650598	4096	Jun	3	2015	samtools
drwxr-xr-x	8	a1650598	4096	Jun	9	2015	samtools-0.1.19
drwxr-xr-x	4	a1650598	4096	Jun	2	2015	samtools-1.2
drwxr-xr-x	5	a1650598	4096	Jun	2	2015	samtools_1_1
drwxr-xr-x	3	a1650598	4096	Jun	2	2015	seqbility-20091110
drwxr-sr-x	4	a1650598	4096	Jan	15	09:46	simwreck
drwxr-sr-x	5	a1650598	4096	Jan	28	13:57	snpEff
drwxr-xr-x	6	a1650598	4096	Jul	1	2015	snpEff_4.1g
drwxr-sr-x	5	a1650598	4096	Jan	28	14:06	snpEff_4.2g
drwxr-xr-x	8	a1650598	4096	Nov	24	09:54	swDMR-1.0.6
drwxr-xr-x	4	a1650598	4096	Jun	2	2015	tmux-1.9a
drwxr-xr-x	4	a1650598	4096	Jul	29	2015	treemix-20150729-git361fac8d
drwxr-xr-x	2	a1650598	4096	Jun	2	2015	trim_galore_zip
drwxr-xr-x	4	a1650598	4096	Jun	2	2015	vcftools-0.1.12b

Within each should have executibles that you will need. For example:



```

[a1650598@phoenix11 Programs]$ ll samtools-1.2/bin/
total 3528
-rwxr-xr-x 1 a1650598 71970 Jun 2 2015 ace2sam
-rwxr-xr-x 1 a1650598 5328 Jun 2 2015 blast2sam.pl
-rwxr-xr-x 1 a1650598 3364 Jun 2 2015 bowtie2sam.pl
-rwxr-xr-x 1 a1650598 18032 Jun 2 2015 export2sam.pl
-rwxr-xr-x 1 a1650598 6125 Jun 2 2015 interpolate_sam.pl
-rwxr-xr-x 1 a1650598 21100 Jun 2 2015 maq2sam-long
-rwxr-xr-x 1 a1650598 21020 Jun 2 2015 maq2sam-short
-rwxr-xr-x 1 a1650598 32675 Jun 2 2015 md5fa
-rwxr-xr-x 1 a1650598 21469 Jun 2 2015 md5sum-lite
-rwxr-xr-x 1 a1650598 6773 Jun 2 2015 novo2sam.pl
-rwxr-xr-x 1 a1650598 48036 Jun 2 2015 plot-bamstats
-rwxr-xr-x 1 a1650598 3208 Jun 2 2015 psl2sam.pl
-rwxr-xr-x 1 a1650598 8961 Jun 2 2015 sam2vcf.pl
-rwxr-xr-x 1 a1650598 3203836 Jun 2 2015 samtools
-rwxr-xr-x 1 a1650598 16478 Jun 2 2015 samtools.pl
-rwxr-xr-x 1 a1650598 11105 Jun 2 2015 seq_cache_populate.pl
-rwxr-xr-x 1 a1650598 3909 Jun 2 2015 soap2sam.pl
-rwxr-xr-x 1 a1650598 7888 Jun 2 2015 varfilter.py
-rwxr-xr-x 1 a1650598 51558 Jun 2 2015 wgsim
-rwxr-xr-x 1 a1650598 4157 Jun 2 2015 wgsim_eval.pl
-rwxr-xr-x 1 a1650598 3583 Jun 2 2015 zoom2sam.pl

[a1650598@phoenix11 bin]$ ./samtools
./samtools: /lib64/libz.so.1: version `ZLIB_1.2.3.3' not found (required by
./samtools)
[a1650598@phoenix11 bin]$ module load zlib
[a1650598@phoenix11 bin]$ ./samtools

Program: samtools (Tools for alignments in the SAM format)
Version: 1.2 (using htlib 1.2.1)

Usage:  samtools <command> [options]

Commands:
  -- indexing

```

Now you can add some of these programs to your PATH variable that is set within bash. You will need to edit your ~/.bashrc file:

This is the start of my bashrc file:

```
[a1650598@afw Refs]$ cat ~/.bashrc

# .bashrc

# Source global definitions
if [ -f /etc/bashrc ]; then
    . /etc/bashrc
fi

# User specific aliases and functions
export PATH="/home/a1650598/.local/bin:\
/localscratch/Programs/bedops_bin:\
/localscratch/Programs/PileOMeth:\
/home/a1650598/acad_ngs_pipeline/bash_bin:\
/home/a1650598/perl5/bin:\
/home/a1650598/bin:\
/home/a1650598/local/HiC-Pro_2.7.3b/bin:\
/localscratch/Programs/bbmap:\
/localscratch/Programs/samtools-1.2/bin:\
/localscratch/Programs/plinkseq-0.10:\
/localscratch/Programs/plink-1.90:\
/localscratch/Programs/gammytools:\
/localscratch/Programs/SailfishBeta-0.9.0_DebianSqueeze/bin:\
/localscratch/Programs/hisat2-2.0.1-beta:\
/localscratch/Programs/SalmonBeta-0.6.1_DebianSqueeze/bin:\
/localscratch/Programs/HiCExplorer:\
/localscratch/Programs/HiCPlotter:\
/localscratch/Programs/segemehl:\
/localscratch/Programs/homer/./bin":$PATH
```

## Reference Sequences and/or Genomes

Obviously, the storage on the cluster gets full fairly quickly. So it makes sense to have a particular directory where you're reference sequences live. This way, multiple users can use the same sequences for their analysis, without having to download or transfer everything all the time

**The directory of this will probably change**

```

[a1650598@afw ~]$ cd /localscratch/Refs/
[a1650598@afw Refs]$ ll
total 7876
drwxrwxr-x.  3 a1650598 afw_users 12288 Jun  1  2015 blast_db
drwxrwxr-x. 10 a1650598 afw_users  4096 Mar 27  2014 Cereals
drwxrwsr-x.  2 a1650598 afw_users    47 May 28  2015 Cyanobacteria
-rw-rw-r--.  1 a1650598 afw_users 25922 May  8  2015 ERCC.fa.gz
drwxrwxr-x. 10 a1650598 afw_users  4096 Mar 24  2014 five_apes
drwxrwxr-x. 10 a1650598 afw_users  4096 Aug 25  2015 fungal_genomes
drwxrwsr-x. 11 a1650598 afw_users  4096 Jan 31 10:30 human
-rw-rw-r--.  1 a1118972 afw_users    47 Nov  2 10:00 info.txt
drwxrwsr-x.  4 a1650598 afw_users  4096 Feb 18 14:46 mouse
drwxrwsr-x.  2 a1650598 a1650598  4096 May 27  2015 PhiX
drwxrwxr-x. 45 a1650598 afw_users  4096 Nov 17 11:08 PlantGenomes
-rw-r-----.  1 a1650598 afw_users 7984536 Sep  3  2015
uchime_reference_dataset_11_03_2015.zip

```